Shared by every arm

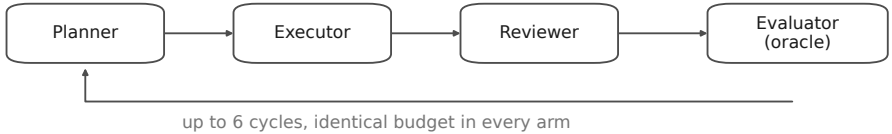


What differs

| | | | measured | verdict |
|---|---|---|---|---|
| A | roles asserted (rhetorical) | *role prompts, in prose* | first pass 0.31–0.80<br>recovery 0.014–0.324 · collapse 0.05 | organized;<br>repairs (1 of 3) |
| B | no role info in any prompt | *one generic prompt for all 3* | first pass 0.31–0.80<br>recovery 0.015–0.247 · collapse 0.91 | same answers;<br>repair falls |
| C | roles enforced (architectural) | *planner schema: no answer field* | first pass 0.28–0.78<br>recovery 0.139–0.395 · attrition 8–17% | organized,<br>but brittle |

Ranges span the three actors. ‘collapse’ = mechanical functional-collapse rate — no judge, because per FINDING-25 the adjudication is not condition-blind.