Shared by every arm

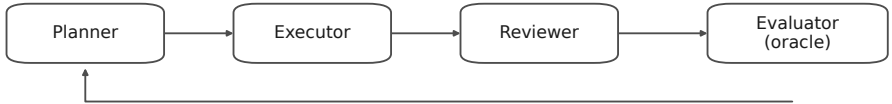

same topology,
routing and budget

What differs

			measured	verdict
A	roles asserted (rhetorical)	*role prompts, in prose*	first pass 0.31–0.80 recovery 0.014–0.324 · collapse 0.05	organized, and repairs
B	no role info in any prompt	*one generic prompt for all 3*	first pass 0.32–0.80 recovery 0.015–0.247 · collapse 0.91	same answers, no repair
C	roles enforced (architectural)	*planner schema: no answer field*	first pass 0.30–0.82 recovery 0.139–0.395 · attrition 8–17%	organized, but brittle

Ranges span the three actors. ‘collapse’ = mechanical functional-collapse rate — no judge, because per FINDING-25 the adjudication is not condition-blind.